

RESEARCH NOTES / EPISODE 01

The question I asked before I knew how to ask it

How my A-level cybersecurity EPQ became the starting point for products built around safer decisions

Presenter: Nathan Brown-Bennett

Target running time: 10–12 minutes

Production script · July 2026

This script is designed for a direct-to-camera research video with document details, restrained data graphics and a short product bridge. Figures should be rechecked immediately before recording.

Episode promise

By the end, viewers should understand why “is cybersecurity secure?” is the wrong test, what the original A-level project got right and wrong, and how the corrected idea led towards a practical device-provisioning product.

CORE LINE

I do not build security theatre. I build tools that make a safer decision easier to understand, repeat and defend.

00:00–00:35 Cold open

PICTURE: Direct to camera. Cut to the original EPQ title page, with the question highlighted.

When I was doing my A-levels, I wrote my first dissertation around a question that sounded straightforward: is cybersecurity actually secure enough for modern-day use?

The question was imperfect. Cybersecurity is not a lock you can label secure or insecure. But hidden inside it was the question I have kept working on ever since: when technology becomes essential, how do we make protection practical enough for people to use—and strong enough to matter?

This is the first in a series about the research behind my products: what I believed, what the evidence now says, and how a dissertation became something you can actually use.

00:35–01:30 Where the project began

PICTURE: Slow pan across the original contents page, survey charts and bibliography. Keep the 2020 date visible.

The original project was written in 2020. I started with physical locks and the idea of perfect security, then moved through information security, cybersecurity, public awareness, human error, artificial intelligence and regulation.

I also ran a small Google Forms quiz among people around me. It asked about hackers, public Wi-Fi, passwords and the differences between cyber and information security. It was useful to me as a student because the answers showed that familiar words did not produce a shared understanding.

But it was not representative research. The sample was small, it was socially concentrated, and the write-up even contains inconsistent response counts. One question asked for a universal minimum password length, which was not a sound premise. So I keep it as part of the history of the work—not as a statistic about the public.

01:30–02:40 What survived the rewrite

GRAPHIC: Four words appear in sequence: No perfection. People matter. Automate carefully. Security is social.

Four ideas survived. First, perfect security is not an operational goal. A control is not useless because a determined attacker may eventually get around it.

Second, people matter. Security advice that nobody understands or can realistically follow is weak security.

Third, automation is necessary. Modern environments create far more events, devices and dependencies than a person can manually check.

And fourth, cybersecurity is social as well as technical. It touches education, business decisions, public services, product design and the way organisations respond when something goes wrong.

02:40–03:55 What the security fields actually protect

GRAPHIC: Use the four-field comparison from the deck, then highlight the same incident crossing all four boundaries.

Before changing the question, I need to separate four terms that are often used as if they mean the same thing.

Computer security protects an individual computing system: its operating system, software, memory, storage and local data. IT security protects an organisation's technology estate—devices, identities, networks, cloud platforms, applications and backups.

Information security protects information in any form, including paper and spoken information, not only digital files. Cybersecurity focuses on harm through connected digital systems, so it can include technology and information, but also people, services and physical processes reachable through them.

The roles overlap because incidents cross boundaries. A stolen cloud credential can require endpoint evidence, IT access revocation, an information-impact assessment and a coordinated cyber response. The labels should clarify ownership and handoffs—not create silos.

03:55–05:10 The question changes

GRAPHIC: Cross out “Is cybersecurity secure?” Replace it with the revised research question.

The biggest correction is the question itself. Security is not a permanent state. It is a decision about risk: what are we protecting, who might attack it, what could the harm be, and what can we reasonably do?

NIST's current Cybersecurity Framework organises that work into six functions: govern, identify, protect, detect, respond and recover.

That final pair changes the conversation. If your definition only asks whether an attack got through, every incident looks like total failure. If you also ask whether it was detected, contained and recovered from, you can distinguish exposure from catastrophe.

So my revised question is this: can contemporary cybersecurity reduce and govern digital risk—and help us recover—at the pace at which technology and attackers change?

05:10–06:50 What the 2026 evidence says

GRAPHIC: Use one number per frame. Always show the source and measurement period.

The current evidence is serious. In the UK Cyber Security Breaches Survey for 2025 to 2026, 43 per cent of businesses said they had identified a breach or attack in the previous twelve months.

The National Cyber Security Centre handled 204 nationally significant incidents in the year to August 2025—more than double the previous year's 89.

But here is the part that matters for the argument: only 30 per cent of businesses reported conducting a cyber-risk assessment, and only 25 per cent had a formal incident-response plan.

Those figures do not prove that cybersecurity does not work. They show that proven practice is unevenly adopted. Some organisations are facing a modern threat environment with no tested plan for the moment prevention fails.

And incident statistics have limits. An organisation with better monitoring may find more attacks. Some incidents are never detected or reported. The number matters—but it is not a simple league table for whether defence works.

06:50–08:10 I would no longer blame the user

PICTURE: Show a hostile-looking security prompt, then simplify it into a clear, safe default.

My original dissertation talked about human error as though people were mainly a vulnerability that automation might replace. I would not frame it that way now.

People do make mistakes. But they make them inside systems that somebody designed. If one ordinary click can expose an entire organisation, we should ask about permissions, authentication, isolation, reporting culture and recovery—not only about the person who clicked.

The better design question is: can we make the safe action easier? Can we use phishing-resistant authentication? Can we remove default passwords? Can somebody report a mistake early without worrying that honesty will punish them?

Training still matters. But training is not a patch for an unsafe product.

08:10–09:25 AI does not remove accountability

GRAPHIC: Split screen: AI-assisted detection on one side; AI-enabled phishing and model risk on the other.

The original project also asked whether artificial intelligence could replace security staff. The better answer is that AI changes both sides of the problem.

It can help defenders triage alerts, inspect code, summarise threat information and spot unusual patterns. It can also help attackers produce convincing messages, automate reconnaissance and work at greater scale.

AI systems introduce their own dependencies and failure modes. Their data, permissions, integrations and outputs need to be governed and tested.

So the useful model is not person versus machine. It is accountable people using automation, with evidence, audit trails and a way to challenge the result.

09:25–11:10 How the research became a product

PICTURE: Show the undergraduate dissertation title, then current Device Provisioning Toolkit discovery and security-guide screens.

The line from this first project to my later work is clearer in hindsight. The EPQ asked why protection is difficult to understand and apply.

My undergraduate dissertation—“How do we decide what we provide?”—made that question specific. When an organisation buys or deploys a connected device, how does it compare the evidence? How does it avoid insecure defaults? How does advice become a repeatable workflow?

That work led to the Device Provisioning Toolkit. It is a research prototype that discovers devices, compares relevant security signals and produces a tailored hardening guide.

It does not promise that a device is perfectly safe. That would repeat the mistake in my first question. It aims to make the decision more informed, the setup more consistent and the reasoning easier to audit.

That is the product principle I have carried forward: do not merely tell people to be more secure. Build what helps them take the next safer action.

11:10–12:00 Answer and close

PICTURE: Return to camera. End on a simple research lineage graphic and the next episode title.

Can cybersecurity keep pace with modern digital risk? My answer is conditional.

Yes—when it is continuous, governed, designed into the product, usable by real people and prepared for recovery.

No—when it is treated as antivirus, an annual training slide, a checklist nobody tests, or a problem delegated to one technical team.

The question I asked at A-level was not quite the right one. But learning why it was wrong gave me a much better direction.

In the next episode, I'll move into my undergraduate dissertation and show how that broad argument became a working comparison and device-provisioning system.

Production notes

- Tone: reflective and precise, not embarrassed by the early work. The point is visible intellectual development.
- Data graphics: display the denominator, period and source; do not animate unrelated numbers together.
- Original document: label every excerpt “2020 A-level submission”.
- Revised document: label every excerpt “2026 critical edition”.
- Product footage: use real current screens. Avoid implying that the prototype certifies a device or guarantees procurement outcomes.
- Accessibility: add captions, describe charts in speech, avoid fast text transitions, and provide the transcript on the website.

On-screen source list

DSIT & Home Office (2026). [Cyber Security Breaches Survey 2025/2026.](#)

NCSC (2025). [Annual Review 2025.](#)

NIST (2024). [Cybersecurity Framework 2.0.](#)

ENISA (2025). [Threat Landscape 2025.](#)

NCSC (2025). [Software Security Code of Practice.](#)